

"What causes antibiotic resistance?"

Tags

Made by Maria Uzun and Alisa Fedorenko

(wait a while for the pictures to load)

1. Get the data

Create a directory for Project 1 materials, and inside it create a new directory for raw data

```
mkdir Project1
cd Project1/
mkdir raw_data
cd raw_data
```

Download the reference sequence of the parental (unevolved, not resistant to antibiotics) *E. coli* strain (fna and gff)

```
wget https://ftp.ncbi.nlm.nih.gov/genomes/all/GCF/000/005/845/GCF_000005845.2_ASM584v2/GCF_000005845.2_ASM584v2_genomic.fna.gz
wget https://ftp.ncbi.nlm.nih.gov/genomes/all/GCF/000/005/845/GCF_000005845.2_ASM584v2/GCF_000005845.2_ASM584v2_genomic.gff.gz
```

Download raw Illumina sequencing reads from shotgun sequencing of an *E. coli* strain that is resistant to the antibiotic ampicillin:

Go to https://figshare.com/articles/dataset/amp_res_2_fastq.zip/10006541/3

Press the button "Download all" and download raw sequence reads into "raw_data" directory

Unarchive downloaded data to the "Project1" directory

```
sudo apt-get install unzip #optionally
```

```
gzip -d -c GCF_000005845.2_ASM584v2_genomic.fna.gz > ../GCF_000005845.2_ASM584v2_genomic.fna
gzip -d GCF_000005845.2_ASM584v2_genomic.gff.gz -c > ../GCF_000005845.2_ASM584v2_genomic.gff
unzip 10006541.zip
```

2. Inspect raw sequencing data manually

Exit from "raw_data" directory

Open the sequence files and verify that the format is correct

```
zcat amp_res_1.fastq.gz|head -20
```

```
@SRR1363257.37 GWZHISEQ01:153:C1W31ACXX:5:1101:14027:2198 length=101  
GGTTGCAGATTCGCGGTCTGCCAGCGCATCACATCTTGTATGTTACGCCGTGGCCTTAGCAATGCTTGAAAGCGAATCGCCTTTGCCACACG  
+  
@?:=:;DBFADH;CAECE@@E:FFHGAE?C?DE<BFGE<>?>FHE4BFFIIFHIBABEECA83;>>@>@CCDC9@@@CC08<@?@BB@9:CC#####  
@SRR1363257.46 GWZHISEQ01:153:C1W31ACXX:5:1101:19721:2155 length=101  
GTATGAGGTTTTGCTGCATTCTCTGNCGGAATATTAACCTCNTNNNNTTATAGTTCAAAGCAAGTACCTGTCTCTTATACACATCTCCGAGCCCACGAGC  
+  
@@<?=D?D==?<AFGDF+AIHEACH#22<:?E8?? :9??GG#0####000;CF=C)4.==CA@@@)=7?C7?E37;3@>;(;>AB#####  
@SRR1363257.77 GWZHISEQ01:153:C1W31ACXX:5:1101:5069:2307 length=101  
GCTTCTCTTAAC TGAGGTCACCATCATGCCGTTAAGTCCCTACCTCTCTTTGCCGGTAACTGTTCCGCCG GATTGCCTTTTATCTGTCTCTTATAACC  
+  
??<DBD;4C2=<BB>:AC;<CF<CE@FE9@E1C@891CD*9?:?3D@DD4?D<DD:0;@A=AEIDDA#####  
@SRR1363257.78 GWZHISEQ01:153:C1W31ACXX:5:1101:5178:2440 length=101  
GCATAAGGACGATCGCTCCAGAGTAAAATAAACGCGCATGTGATACTACAATACCAATGGTGAAGTTACGGGACTTAAACAAACTGAGATCAAGAATC  
+  
CCCCFFFFHHHHJJJJJJJJJJJFFHIJJJJJJJJJJJJJJJJJJJJJIHHHHHFDEF;AEEDDDDBBACDDDCDDDDCCDDDDCCDC3  
@SRR1363257.96 GWZHISEQ01:153:C1W31ACXX:5:1101:6707:2460 length=101  
TCATTAAGCCGTGGTGGATGTGCCATAGCGCACCGCAAAGTTAAGAAACCGAATATTGGGTTAGTCTTGTTCATAATTGTTGCAATGAAACGCGGTGAA  
+  
CCCCFFFFHHHHJJIIIIIJJJJJJJJGIJJJJJJJJIIIGHJJJJJJIIJDHFFFFFFEDACDDDCDDDDCCDDECACCCDACDDDDCCDDDDBD@A
```

```
zcat amp_res_2.fastq.gz|head -20
```

```
@SRR1363257.37 GWZHISEQ01:153:C1W31ACXX:5:1101:14027:2198 length=101  
GATCTAAGCTGAAGCCAGGCCAAAGTTTGACGATTGGTGCAGGCAGTAGCGCACAGCGACTGGCAAACAACAGCGATAGCATTACGTATCGTGTGCGCAAA  
+  
???BDB:DFHBF@?; ; +A; AFGH; ABHFHHGE@?: B: ?? @D> @; F? D8<<F8AA9EHHD8' . . ; 5?A?A992(' , (59CC3@C>22::A238+>B<>B<  
@SRR1363257.46 GWZHISEQ01:153:C1W31ACXX:5:1101:19721:2155 length=101  
GTACTTGCTTTGNACTATAATATGCACGGAGNTAATATTCGCTCAGAGAATGCAGCAAAACCTCATACCTGTCTCTTATACACATCTGACGCTGCCGACGA  
+  
; @DB?B; CFBB#2<; CB:FHC@:<A?C:: #1:86:BG9:8?8688?888EBF; 783)=6-7=CC; ECD); ?7; >>AE; >(5;->AC@; B@; 8?####  
@SRR1363257.77 GWZHISEQ01:153:C1W31ACXX:5:1101:5069:2307 length=101  
ATAATAGGCAATCGCGTCGGAACAGTTACCGGCCAAAGAGAGGCGAGGACTTAACGGCATGATGGTGCACCTCAGTTAAGAGAAGCCTGTCTCTTATACACA  
+  
+=?; :2,+ ++ AC:C:2@F6: CD:B09B?4)8@'( '=))8 =; =( (5=4@?; @6; @?@BB; (535::>:>3(: (: (44:@:: @3((9<32+: :@(4@4+:> C3  
@SRR1363257.78 GWZHISEQ01:153:C1W31ACXX:5:1101:5178:2440 length=101  
ATATTAACAGTAGTATCAGTTATTTCTCTGATCTCTTTAGTCATTTGGGAGTCGACCTCAGAGAACC CGATTCTTGATCTCAGTTT GTTTAAGTCCC GTAA  
+  
BCCFFFFFH HHHIJJJIIJJJJJJJJJJJIJJJJJJJJJ JIHJJJJJJ IIGGGHIJJJJJJJJJJJJGH HHHFF FFFFEEEFEEED?AACDCDD DB  
@SRR1363257.96 GWZHISEQ01:153:C1W31ACXX:5:1101:6707:2460 length=101  
GTTTTACCGCGTTTCATTGCAACAATTATGAACAAGACTAAACCCAATATTCG G TTTCTTA ACTTTGCGGTGCGCTATGGCACATCCACCACG G CT TAAT  
+  
CCCCFFFF FH HH IJJJJJJJJJJJJJJJJJJJJJJJJJJJJJJ JJFI IIJJFIGJJJJGIHHHHHG FFDDDDDDDDDDDDDDDDDABDDDDDDDCD
```

```
zcat GCF_000005845.2_ASM584v2_genomic.fna.gz|head -20
```

>NC_000913.3 *Escherichia coli* str. K-12 substr. MG1655, complete genome
AGCTTTTCTATTCTGACTGCAACGGGCAATATGTCTCTGTGTGGATTAAAAAAGAGTGTCTGATAGCAGCTTCTGAAGT
GTTACCTGCCGTGAGTAAATTAATAATTTTATTGACTTAGGTCATAAATACTTTAACCAATATAGGCATAGCGCACAGAC
AGATAAAAAATTACAGGTACACAACATCCATGAAACGCATTAGCACCACCATTACCACCACCATCACCATTACCACAGGT
AACGGTGCGGGCTGACGCGTACAGGAAACACAGAAAAAGCCCGACCTGACAGTGCGGGCTTTTTTTTTTCGACCAAAGG
TAACGAGGTAACAACCATGCGAGTGTTGAAGTTCGGCGGTACATCAGTGGCAAATGCAGAACGTTTTCTGCGTGTTGCCG
ATATTCTGGAAGCAATGCCAGGCGAGGGGAGGTGGCCACCGTCTCTCTGCCCCGCCAAAATCACCAACCACCTGGTG
GCGATGATTGAAAAAACCATTAGCGGCCAGGATGCTTTACCAATATCAGCGATGCCGAACGATTTTTTGCCGAACTTTT
GACGGGACTCGCCGCCGCCAGCCGGGTTCCCGCTGGCGCAATTGAAAACCTTCGTCGATCAGGAATTTGCCCAAATAA
AACATGTCCTGCATGGCATTAGTTTGTTGGGGCAGTGCCCGGATAGCATCAACGCTGCGCTGATTTGCCGTGGCGAGAAA
ATGTCGATCGCCATTATGGCCGGCGTATTAGAAGCGCGCGGTACACAACGTTACTGTTATCGATCCGGTCGAAAAACTGCT
GGCAGTGGGGCATTACCTCGAATCTACCGTCGATATTGCTGAGTCCACCCGCCGATTGCGGCAAGCCGCATTCCGGCTG
ATCACATGGTGCTGATGGCAGTTTTACCGCCGGTAATGAAAAAGGCGAACTGGTGCTGTTGGACGCAACGTTCCGAC
TACTCTGCTGCGGTGCTGGCTGCCTGTTTACGCGCCGATTGTTGCGAGATTTGGACGGACGTTGACGGGGTCTATACCTG
CGACCCCGCTCAGGTGCCCGATGCGAGGTTGTTGAAGTCGATGTCCTACCAGGAAGCGATGGAGCTTTCCTACTTCGGCG
CTAAAGTTCTTCACCCCCGCACCATTACCCCATCGCCAGTTCAGATCCCTTGCTGATTAAAAATACCGGAAATCCT
CAAGCACCAGGTACGCTCATTGGTGCCAGCCGTGATGAAGACGAATTACCGGTCAAGGGCATTTCGAATACAT
GGCAATGTTACGCGTTTCTGGTCCGGGGATGAAAGGGATGGTCGGCATGGCGGCGCGCTCTTGCAGCGATGTCACGCG
CCCGTATTTCCGTGGTGCTGATTACGCAATCATCTTCGAATACAGCATCAGTTTCTGCGTTCCACAAAGCGACTGTGTG
CGAGCTGAACGGGCAATGCAGGAAGAGTTCTACCTGGAACGAAAGAAGGCTTACTGGAGCCGCTGGCAGTGACGGAACG

```
zcat GCF_000005845.2_ASM584v2_genomic.gff.gz|head -20
```

```

#lgff-version 3
#lgff-spec-version 1.21
#|processor NCBI annotwriter
#|genome-build ASM584v2
#|genome-build-accession NCBI_Assembly:GCF_000005845.2
##sequence-region NC_000913.3 1 4641652
##species https://www.ncbi.nlm.nih.gov/Taxonomy/Browser/wwwtax.cgi?id=511145
NC_000913.3 RefSeq region 1 4641652 . + . ID=NC_000913.3:1..4641652;Dbxref=taxon:511145;Is_circular=true;Name=ANONYMOUS;gbkey=Src;genome=chromosome;mol_type=genomic DNA;strain=K-12;
substrain=MG1655
NC_000913.3 RefSeq gene 190 255 . + . ID=gene-b0001;Dbxref=ASAP:ABE-0000006,ECOCYC:EG11277,GenID:944742;Name=thrL;gbkey=Gene;gene=thrL;gene_biotype=protein_coding;gene_synonym=
ECK0001;locus_tag=b0001
NC_000913.3 RefSeq CDS 190 255 . + 0 ID=cds-NP_414542.1;Parent=gene-b0001;Dbxref=UniProtKB/Swiss-Prot:P0AD86,GenBank:NP_414542.1,ASAP:ABE-0000006,ECOCYC:EG11277,GenID:944742;Name=NP_414542.1;gbkey=CDS;gene=thrL;locus_tag=b0001;orig_transcript_id=gnl|b0001|mrna.NP_414542;product=thr operon leader peptide;protein_id=NP_414542.1;transl_table=11
NC_000913.3 RefSeq gene 337 2799 . + . ID=gene-b0002;Dbxref=ASAP:ABE-0000008,ECOCYC:EG10998,GenID:945803;Name=thrA;gbkey=Gene;gene=thrA;gene_biotype=protein_coding;gene_synonym=
ECK0002;Hs_thrA1,thrA2,thrD;locus_tag=b0002
NC_000913.3 RefSeq CDS 337 2799 . + 0 ID=cds-NP_414543.1;Parent=gene-b0002;Dbxref=UniProtKB/Swiss-Prot:P00561,GenBank:NP_414543.1,ASAP:ABE-0000008,ECOCYC:EG10998,GenID:945803;Name=NP_414543.1;gbkey=CDS;gene=thrA;locus_tag=b0002;orig_transcript_id=gnl|b0002|mrna.NP_414543;product=fused aspartate kinase/homoserine dehydrogenase 1;protein_id=NP_414543.1;transl_table=11
NC_000913.3 RefSeq gene 2801 3733 . + . ID=gene-b0003;Dbxref=ASAP:ABE-0000010,ECOCYC:EG10999,GenID:947498;Name=thrB;gbkey=Gene;gene=thrB;gene_biotype=protein_coding;gene_synonym=
ECK0003;locus_tag=b0003
NC_000913.3 RefSeq CDS 2801 3733 . + 0 ID=cds-NP_414544.1;Parent=gene-b0003;Dbxref=UniProtKB/Swiss-Prot:P00547,GenBank:NP_414544.1,ASAP:ABE-0000010,ECOCYC:EG10999,GenID:947498;Name=NP_414544.1;gbkey=CDS;gene=thrB;locus_tag=b0003;orig_transcript_id=gnl|b0003|mrna.NP_414544;product=homoserine kinase;protein_id=NP_414544.1;transl_table=11
NC_000913.3 RefSeq gene 3734 5020 . + . ID=gene-b0004;Dbxref=ASAP:ABE-0000012,ECOCYC:EG11000,GenID:945198;Name=thrC;gbkey=Gene;gene=thrC;gene_biotype=protein_coding;gene_synonym=
ECK0004;locus_tag=b0004
NC_000913.3 RefSeq CDS 3734 5020 . + 0 ID=cds-NP_414545.1;Parent=gene-b0004;Dbxref=UniProtKB/Swiss-Prot:P00934,GenBank:NP_414545.1,ASAP:ABE-0000012,ECOCYC:EG11000,GenID:945198;Name=NP_414545.1;gbkey=CDS;gene=thrC;locus_tag=b0004;orig_transcript_id=gnl|b0004|mrna.NP_414545;product=threonine synthase;protein_id=NP_414545.1;transl_table=11
NC_000913.3 RefSeq gene 5234 5530 . + . ID=gene-b0005;Dbxref=ASAP:ABE-0000015,ECOCYC:G6081,GenID:944747;Name=yaaX;gbkey=Gene;gene=yaaX;gene_biotype=protein_coding;gene_synonym=EC
K0005;locus_tag=b0005
NC_000913.3 RefSeq CDS 5234 5530 . + 0 ID=cds-NP_414546.1;Parent=gene-b0005;Dbxref=UniProtKB/Swiss-Prot:P75616,GenBank:NP_414546.1,ASAP:ABE-0000015,ECOCYC:G6081,GenID:944747;Name=NP_414546.1;gbkey=CDS;gene=yaaX;locus_tag=b0005;orig_transcript_id=gnl|b0005|mrna.NP_414546;product=DUF2502 domain-containing protein YaaX;protein_id=NP_414546.1;transl_table=11
NC_000913.3 RefSeq gene 5683 6459 . - . ID=gene-b0006;Dbxref=ASAP:ABE-0000018,ECOCYC:EG10011,GenID:944749;Name=yaaA;gbkey=Gene;gene=yaaA;gene_biotype=protein_coding;gene_synonym=
ECK0006;locus_tag=b0006
NC_000913.3 RefSeq CDS 5683 6459 . - 0 ID=cds-NP_414547.1;Parent=gene-b0006;Dbxref=UniProtKB/Swiss-Prot:P0AB13,GenBank:NP_414547.1,ASAP:ABE-0000018,ECOCYC:EG10011,GenID:944749;Name=NP_414547.1;gbkey=CDS;gene=yaaA;locus_tag=b0006;orig_transcript_id=gnl|b0006|mrna.NP_414547;product=DNA binding and peroxide stress response protein YaaA;protein_id=NP_414547.1;transl_table=11

```

Open the entire fasta reference file. Do you notice anything different about it? I see circular chromosome sequence

```
zcat GCF_000005845.2_ASM584v2_genomic.fna.gz
```

How many reads are in each fastq file?

```
conda install -c bioconda seqkit
seqkit stats amp_res_1.fastq.gz
```

File	Format	Type	Num_seqs	Sum_len	Min_len	Avg_len	Max_len
amp_res_1.fastq.gz	FASTQ	DNA	455,876	46,043,476	101	101	101

```
seqkit stats amp_res_2.fastq.gz
```

File	Format	Type	Num_seqs	Sum_len	Min_len	Avg_len	Max_len
amp_res_2.fastq.gz	FASTQ	DNA	455,876	46,043,476	101	101	101

3. Inspect raw sequencing data with FastQC. Filtering the reads.

```
cd ..
mkdir fastqc
cd fastqc
fastqc -o ./ ../raw_data/amp_res_1.fastq.gz ../raw_data/amp_res_2.fastq.gz
```

[amp_res_1_fastqc.html](#)

[amp_res_2_fastqc.html](#)

Do the basic statistics match what you calculated for the number of reads last time?

Yes, it matches

Basic Statistics

Measure	Value
Filename	amp_res_1.fastq.gz
File type	Conventional base calls
Encoding	Sanger / Illumina 1.9
Total Sequences	455876
Total Bases	46 Mbp
Sequences flagged as poor quality	0
Sequence length	101
%GC	50

In generated files, we see red circles at **Per base sequence quality** and at **Per tile sequence quality**


Summary

- ✓ [Basic Statistics](#)
- ✗ [Per base sequence quality](#)
- ✗ [Per tile sequence quality](#)
- ✓ [Per sequence quality scores](#)
- ! [Per base sequence content](#)
- ! [Per sequence GC content](#)
- ✓ [Per base N content](#)
- ✓ [Sequence Length Distribution](#)
- ✓ [Sequence Duplication Levels](#)
- ✓ [Overrepresented sequences](#)
- ✓ [Adapter Content](#)

What do they mean?

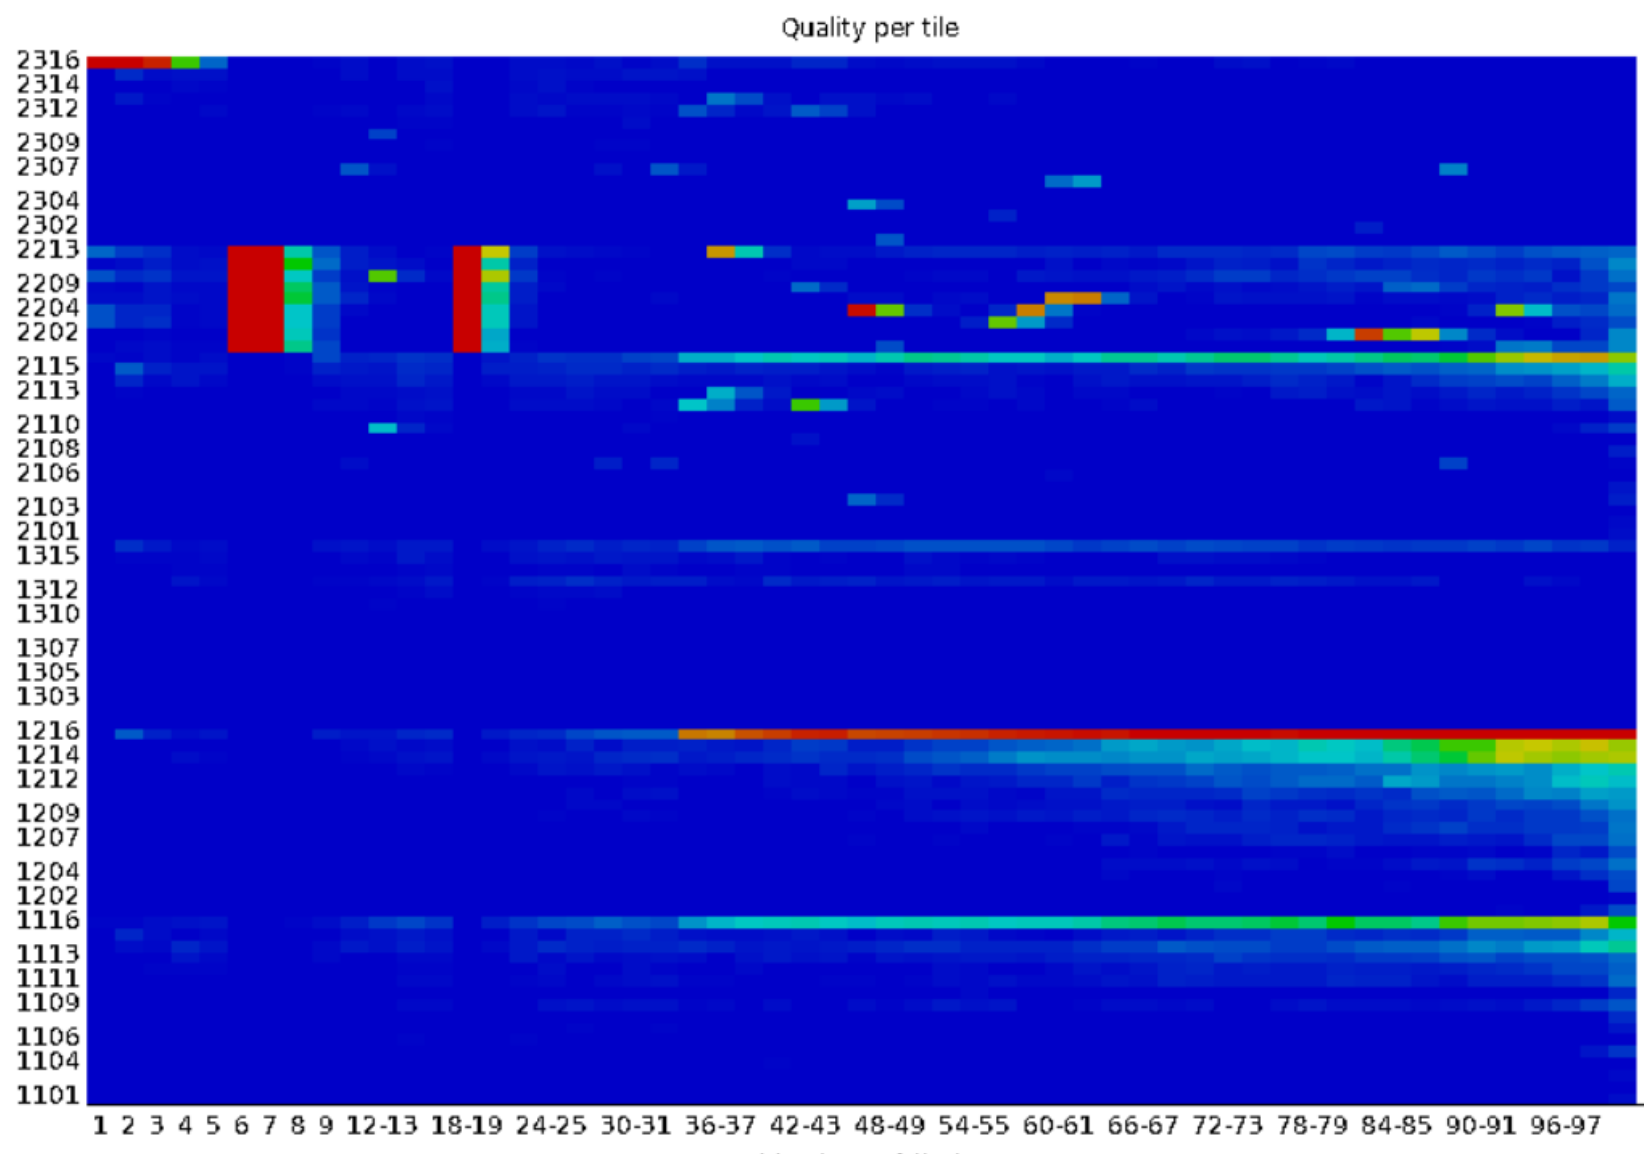
Red bars in the "Per Base Sequence Quality" section usually indicate that in our sequences there are reads where the quality scores drop significantly. These quality drops may suggest that there are regions in sequences with lower accuracy, potentially due to sequencing errors, adapter contamination, or other issues.

✗ Per base sequence quality



Red circles in the "Per Tile Sequence Quality" section indicate lower quality or unusual patterns on individual tiles of the flow cell (e.g. air bubbles or reagent problems).

✖ Per tile sequence quality



What we should do about anything FastQC identified as unusual?

We should explore them accurately, and, when possible, eliminate them.

4. Filtering the reads

Install Trimmomatic using conda.

```
conda install -c bioconda trimmomatic
cd ..
mkdir trimmomatic
cd trimmomatic
```

Trimmomatic was run in paired-end mode, with the following parameters:

- Cut bases off the start of a read if quality is below 20
- Cut bases off the end of a read if quality is below 20
- Trim reads using a sliding window approach, with window size 10 and average quality within the window 20.
- Drop the read if it is below length 20.

```
trimmomatic PE -threads 16 -phred33 ../raw_data/amp_res_1.fastq.gz ../raw_data/amp_res_2.fastq.gz trimmomatic_20_forward_paired.fq.gz trimmomatic_20_forward_unpaired.fq.gz trimmomatic_20_reverse_paired.fq.gz trimmomatic_20_reverse_unpaired.fq.gz LEADING:20 TRAILING:20 SLIDINGWINDOW:10:20 MINLEN:20
```

Trimmomatic reports back stats, on how many paired reads were kept

```
q.gz trimmomatic_20_forward_unpaired.fq.gz trimmomatic_20_reverse_paired.fq.gz trimmomatic_20_reverse_unpaired.fq.gz LEADING:20 TRAILING:20 SLIDINGWINDOW:10:20 MINLEN:20
TrimmomaticPE: Started with arguments:
-threads 16 -phred33 ../raw_data/amp_res_1.fastq.gz ../raw_data/amp_res_2.fastq.gz trimmomatic_20_forward_paired.fq.gz trimmomatic_20_forward_unpaired.fq.gz trimmomatic_20_reverse_paired.fq.gz trimmomatic_20_reverse_unpaired.fq.gz LEADING:20 TRAILING:20 SLIDINGWINDOW:10:20 MINLEN:20
Input Read Pairs: 455876 Both Surviving: 446259 (97.89%) Forward Only Surviving: 9216 (2.02%) Reverse Only Surviving: 273 (0.06%) Dropped: 128 (0.03%)
TrimmomaticPE: Completed successfully
```

Manual checking

```
seqkit stats trimmomatic_20_forward_paired.fq.gz
```

File	Format	Type	Num_seqs	Sum_len	Min_len	Avg_len	Max_len
trimmomatic_20_forward_paired.fq.gz	FASTQ	DNA	446,259	42,003,868	20	94.1	101

```
seqkit stats trimmomatic_20_reverse_paired.fq.gz
```

File	Format	Type	Num_seqs	Sum_len	Min_len	Avg_len	Max_len
trimmomatic_20_reverse_paired.fq.gz	FASTQ	DNA	446,259	41,649,154	20	93.3	101


To see how trimming affected the overall quality of the data, we repeated the fastqc analysis from section 3, but this time on the _1P.fq and _2P.fq files.

```
fastqc trimmomatic_20_forward_paired.fq.gz trimmomatic_20_reverse_paired.fq.gz -o ../fastqc/
```

As a result, we see that according to **Per base sequence quality** everything was cut/trimmed correctly.


✔

Per base sequence quality



Per tile sequence quality was also improved, except for one part, which, apparently, was caused by flowcell problems.

✖ Per tile sequence quality




Next, we increased the quality score at all steps to 30

```
trimmomatic PE -threads 16 -phred33 ../raw_data/amp_res_1.fastq.gz ../raw_data/amp_res_2.fastq.gz trimmomatic_30_forward_paired.fq.gz trimmomatic_30_forward_unpaired.fq.gz trimmomatic_30_reverse_paired.fq.gz trimmomatic_30_reverse_unpaired.fq.gz LEADING:30 TRAILING:30 SLIDINGWINDOW:10:30 MINLEN:30
fastqc trimmomatic_30_forward_paired.fq.gz trimmomatic_30_reverse_paired.fq.gz -o ../fastqc/
```

As a result we see, that all reads are in green zone according to their quality scores, which is great.



Per base sequence quality



5 . Aligning sequences to reference

5.1 Index the reference file

Run bwa index on the reference sequence with the default options

```
cd ..
mkdir alignment
cd alignment
#conda install -c bioconda bwa
cp ../raw_data/GCF_000005845.2_ASM584v2_genomic.fna.gz .
gzip -d -c GCF_000005845.2_ASM584v2_genomic.fna.gz > GCF_000005845.2_ASM584v2_genomic.fna
bwa index GCF_000005845.2_ASM584v2_genomic.fna
ls
```

Creates index files

```
GCF_000005845.2_ASM584v2_genomic.fna.gz      GCF_000005845.2_ASM584v2_genomic.fna.gz.ann  GCF_000005845.2_ASM584v2_genomic.fna.gz.pac
GCF_000005845.2_ASM584v2_genomic.fna.gz.amb  GCF_000005845.2_ASM584v2_genomic.fna.gz.bwt  GCF_000005845.2_ASM584v2_genomic.fna.gz.sa
```

5.2 Aligning reads

```
bwa mem -t 16 GCF_000005845.2_ASM584v2_genomic.fna ../trimmomatic/trimmomatic_20_forward_paired.fq.gz ../trimmomatic/trimmomatic_20_reverse_paired.fq.gz > alignment.sam
```

If nothing happens, press Enter 😊


```

reverse_paired.fq.gz > alignment_tr.sam
[M::bwa_idx_load_from_disk] read 0 ALT contigs
[M::process] read 892518 sequences (83653022 bp)...
[M::mem_pestat] # candidate unique pairs for (FF, FR, RF, RR): (109, 428813, 0, 115)
[M::mem_pestat] analyzing insert size distribution for orientation FF...
[M::mem_pestat] (25, 50, 75) percentile: (71, 117, 175)
[M::mem_pestat] low and high boundaries for computing mean and std.dev: (1, 383)
[M::mem_pestat] mean and std.dev: (121.75, 68.23)
[M::mem_pestat] low and high boundaries for proper pairs: (1, 487)
[M::mem_pestat] analyzing insert size distribution for orientation FR...
[M::mem_pestat] (25, 50, 75) percentile: (143, 183, 228)
[M::mem_pestat] low and high boundaries for computing mean and std.dev: (1, 398)
[M::mem_pestat] mean and std.dev: (187.49, 63.10)
[M::mem_pestat] low and high boundaries for proper pairs: (1, 483)
[M::mem_pestat] skip orientation RF as there are not enough pairs
[M::mem_pestat] analyzing insert size distribution for orientation RR...
[M::mem_pestat] (25, 50, 75) percentile: (82, 126, 220)
[M::mem_pestat] low and high boundaries for computing mean and std.dev: (1, 496)
[M::mem_pestat] mean and std.dev: (130.49, 81.88)
[M::mem_pestat] low and high boundaries for proper pairs: (1, 634)
[M::mem_pestat] skip orientation FF
[M::mem_pestat] skip orientation RR
[M::mem_process_seqs] Processed 892518 reads in 42.850 CPU sec, 4.231 real sec
[main] Version: 0.7.17-r1188
[main] CMD: bwa mem -t 16 GCF_000005845.2_ASM584v2_genomic.fna ../trimmomatic/trimmomatic_20_forward_paired.fq.gz ../trimmomatic/trimmomatic_20_reverse_paired.fq.gz
[main] Real time: 19.631 sec; CPU: 46.488 sec

```

5.3. Compress SAM file

To compress and sort the sam file with the commands below.

A compressed sam file is called a bam file

```
samtools view -S -b alignment.sam > alignment.bam
```

To get some basic statistics:

```
samtools flagstat alignment.bam
```

```

912095 + 0 in total (QC-passed reads + QC-failed reads)
911752 + 0 primary
0 + 0 secondary
343 + 0 supplementary
0 + 0 duplicates
0 + 0 primary duplicates
910940 + 0 mapped (99.87% : N/A)
910597 + 0 primary mapped (99.87% : N/A)
911752 + 0 paired in sequencing
455876 + 0 read1
455876 + 0 read2
907442 + 0 properly paired (99.53% : N/A)
909488 + 0 with itself and mate mapped
1109 + 0 singletons (0.12% : N/A)
0 + 0 with mate mapped to a different chr
0 + 0 with mate mapped to a different chr (mapQ>=5)

```

What percentage of reads are mapped?

99.87%

5.4 Sort and index BAM file

a) Sort bam file by sequence coordinate on reference

```
samtools sort alignment.bam -o alignment_sorted.bam
```

b) Index bam file for faster search:

```
samtools index alignment_sorted.bam
```

6. Variant calling

The solution is to make an intermediate file type called an mpileup, because it goes through each position and “piles up” the reads, tabulating the number of bases that match or don’t match the reference. Mpileup requires a sorted, indexed bam file. For this, run the basic command below.

```

cd ..
mkdir variant_calling
cd variant_calling
samtools mpileup -f ../alignment/GCF_000005845.2_ASM584v2_genomic.fna ../alignment/alignment_sorted.bam > my.mpileup

```

To call actual variants, we will be using a program called VarScan (variant scanner)

Download version 2.4.0 to the variant_calling directory from here <https://github.com/dkoboldt/varscan/blob/master/VarScan.v2.4.0.jar>

Go call help for a command mpileup2snp

```
java -jar VarScan.v2.4.0.jar mpileup2snp -h
```

```
Only SNPs will be reported
Warning: No p-value threshold provided, so p-values will not be calculated
Min coverage: 8
Min reads2: 2
Min var freq: 0.2
Min avg qual: 15
P-value thresh: 0.01
USAGE: java -jar VarScan.jar mpileup2cns [pileup file] OPTIONS
       mpileup file - The SAMtools mpileup file

OPTIONS:
--min-coverage  Minimum read depth at a position to make a call [8]
--min-reads2    Minimum supporting reads at a position to call variants [2]
--min-avg-qual  Minimum base quality at a position to count a read [15]
--min-var-freq  Minimum variant allele frequency threshold [0.01]
--min-freq-for-hom  Minimum frequency to call homozygote [0.75]
--p-value       Default p-value threshold for calling variants [99e-02]
--strand-filter Ignore variants with >90% support on one strand [1]
--output-vcf    If set to 1, outputs in VCF format
--vcf-sample-list  For VCF output, a list of sample names in order, one per line
--variants      Report only variant (SNP/indel) positions [0]
```

For SNPs checking run the following command

```
java -jar VarScan.v2.4.0.jar mpileup2snp my.mpileup --min-var-freq 0.2 --variants --output-vcf 1 > VarScan_results.vcf
```

6 variant positions reported (6 SNP, 0 indel)

```
Only SNPs will be reported
Warning: No p-value threshold provided, so p-values will not be calculated
Min coverage: 8
Min reads2: 2
Min var freq: 0.2
Min avg qual: 15
P-value thresh: 0.01
Reading input from my.mpileup
4641524 bases in pileup file
9 variant positions (6 SNP, 3 indel)
0 were failed by the strand-filter
6 variant positions reported (6 SNP, 0 indel)
```

7. Variant effect prediction

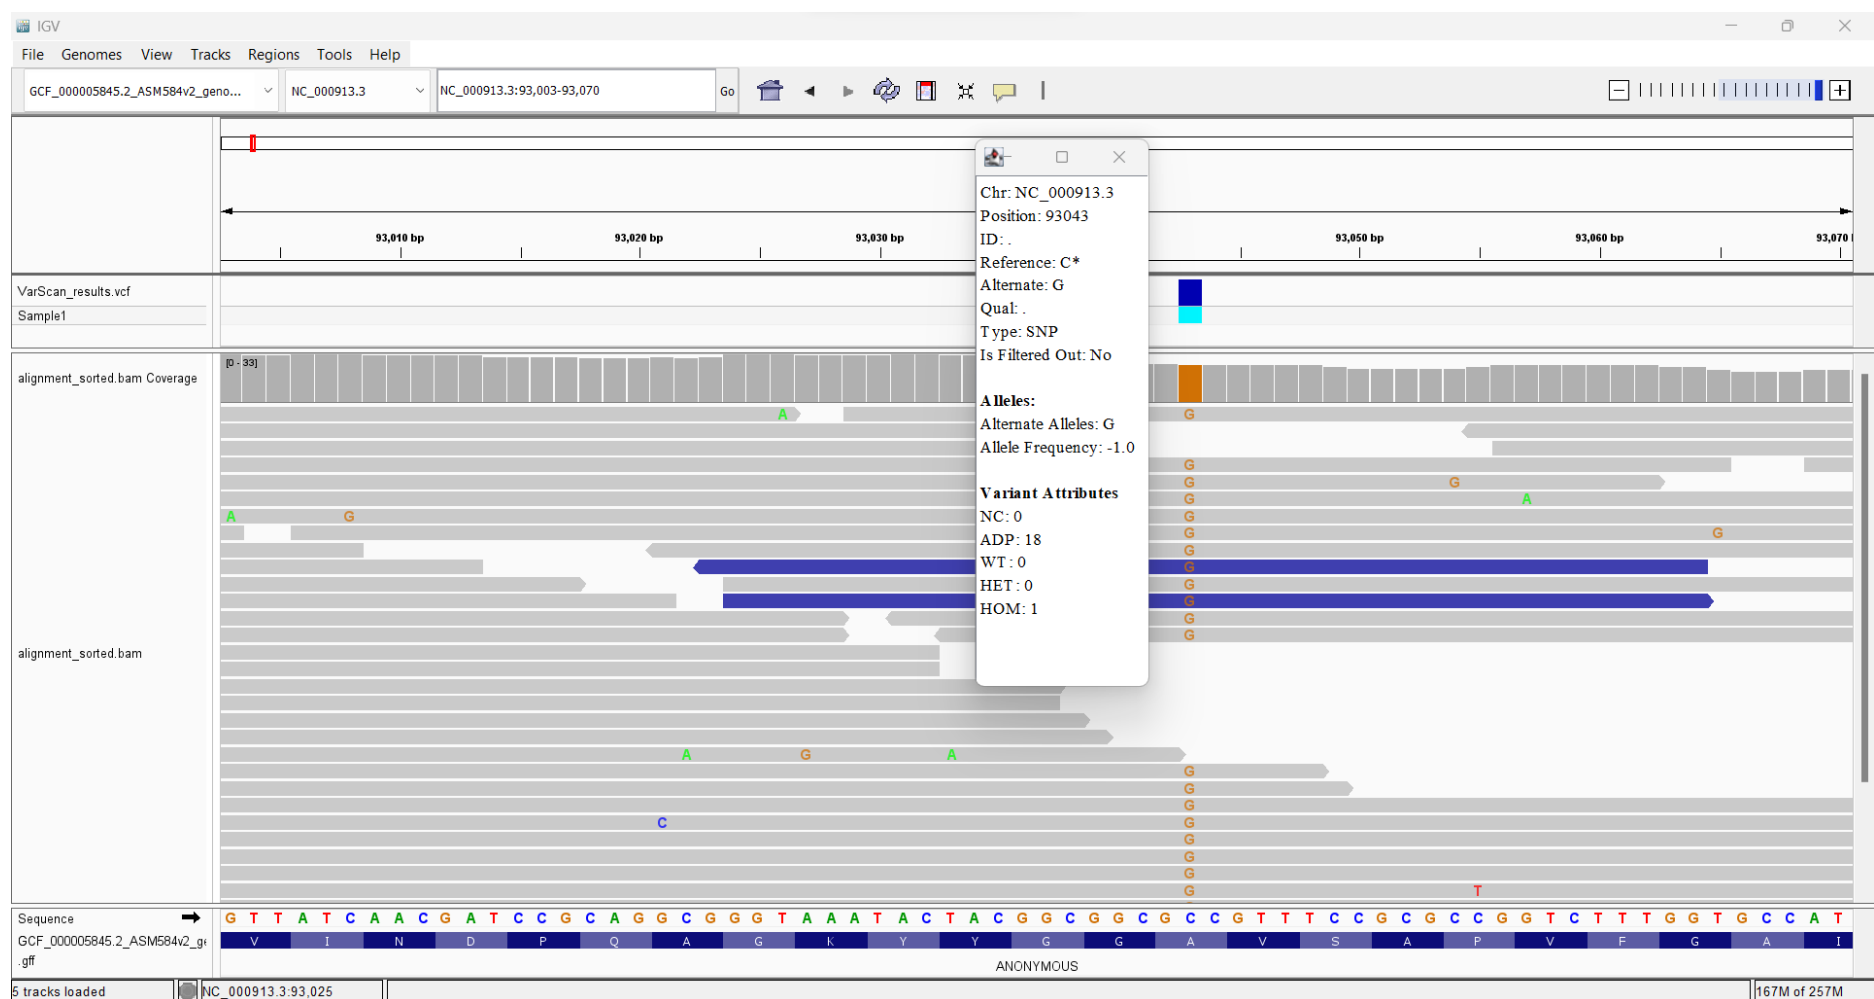
Data visualization in IGV browser:

Select "Genomes", "Load Genome from File" and select our reference genome. Then select "File", "Open from file" and select your BAM file. Add also vcf file and annotation in gff format.

Results of visualization:

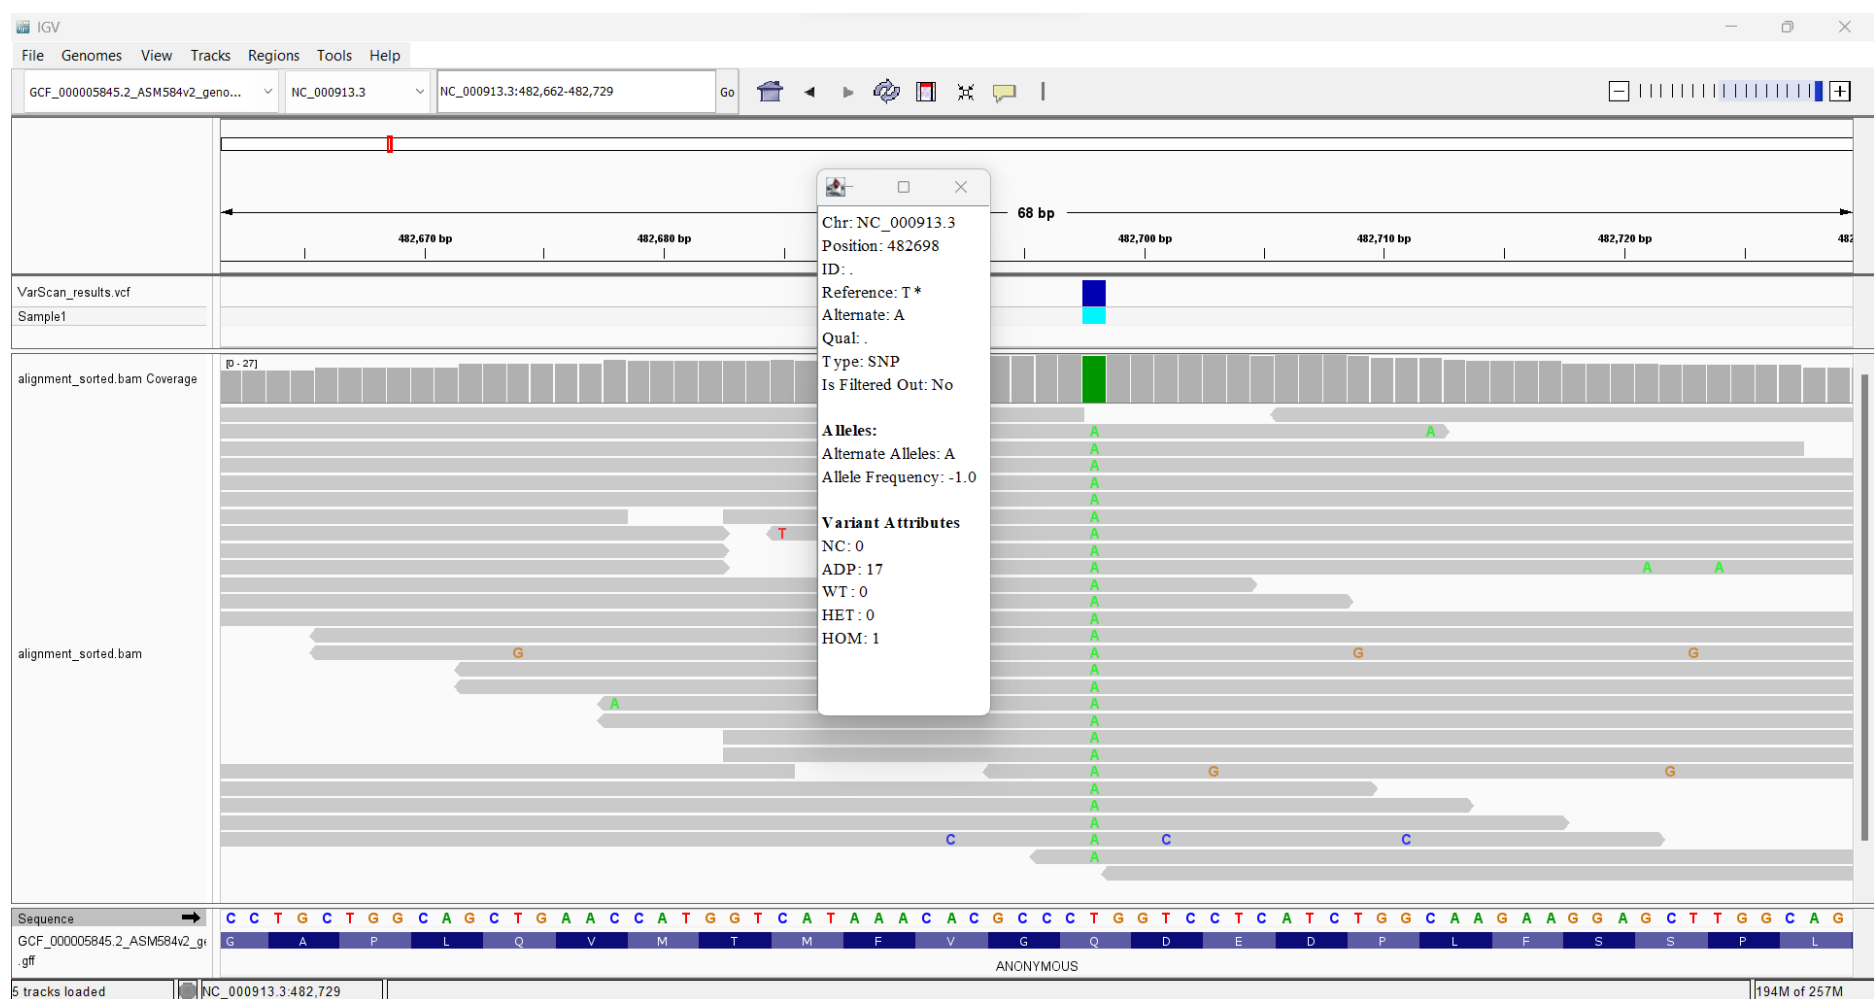
SNP 1

- Position: 93043
- Reference: C
- Alternate: G
- Gene: ftst
- Effect: Missense variant, changing amino acid from Alanine (Ala) to Glycine (Gly)



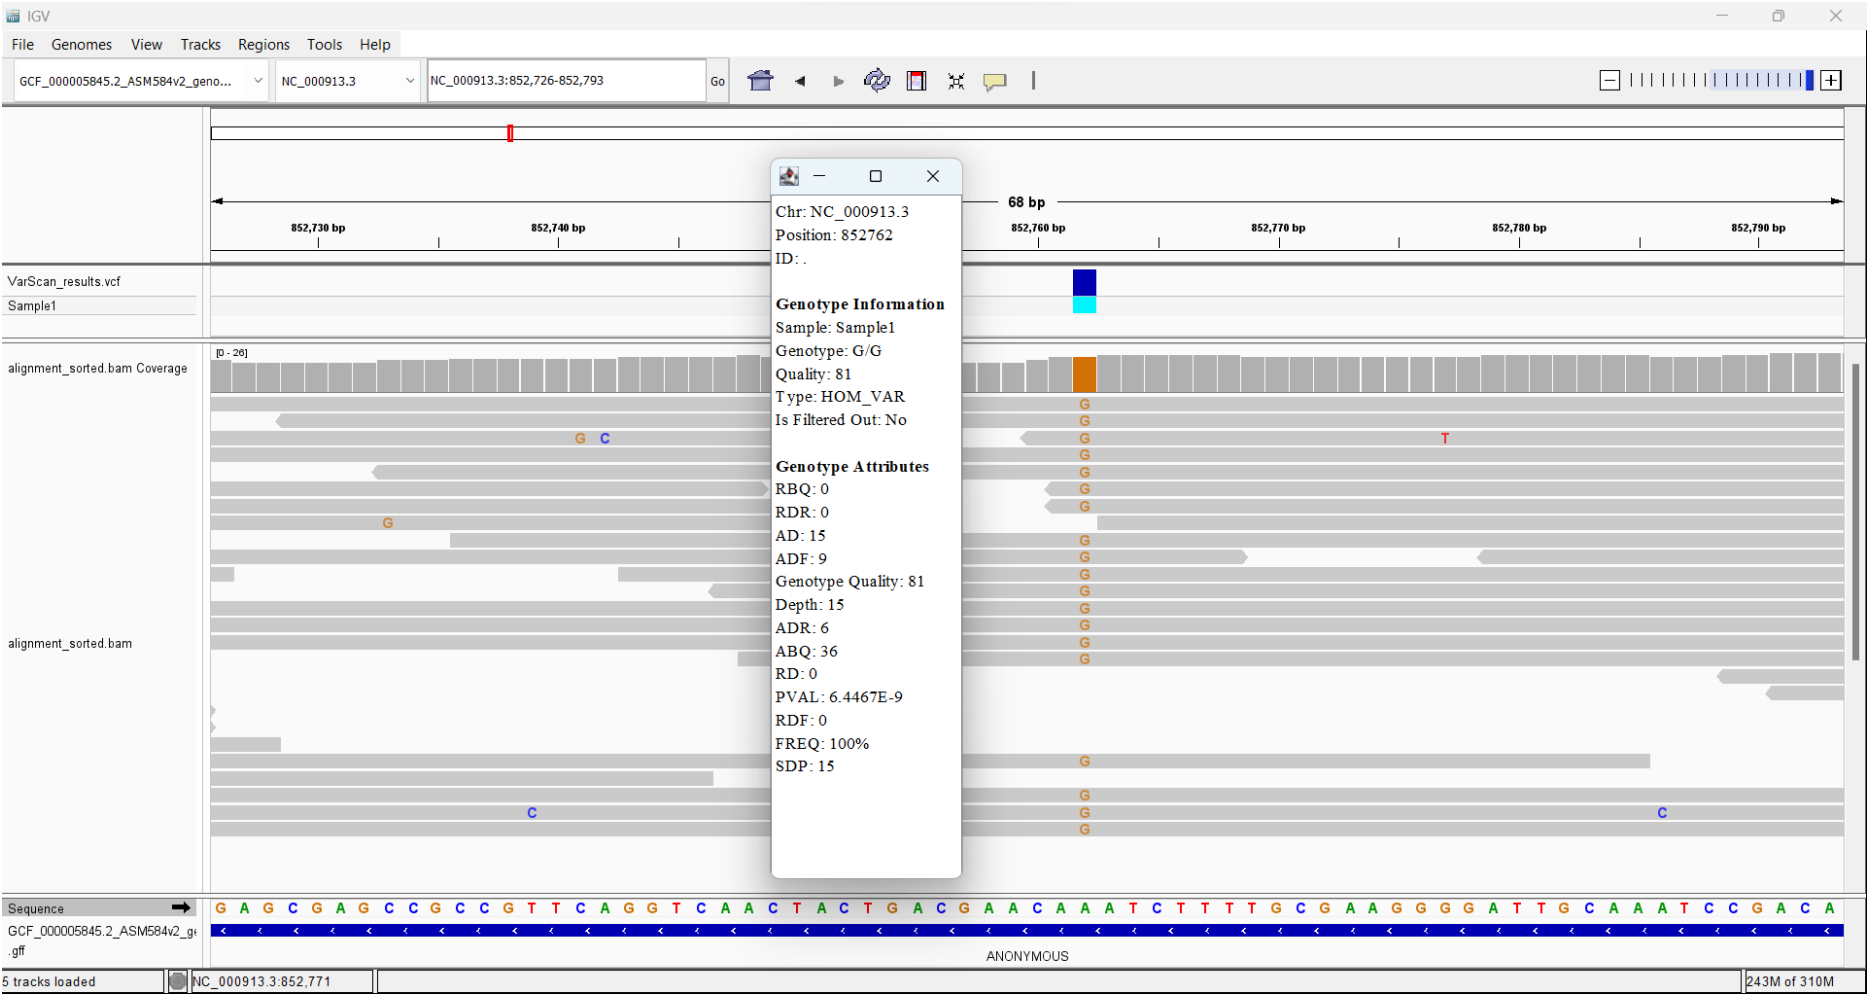
SNP 2

- Position: 482698
- Reference: A
- Alternate: T
- Gene: acrB
- Effect: Missense variant, changing amino acid from Glutamine (Gln/Q) to Leucine (Leu/L)



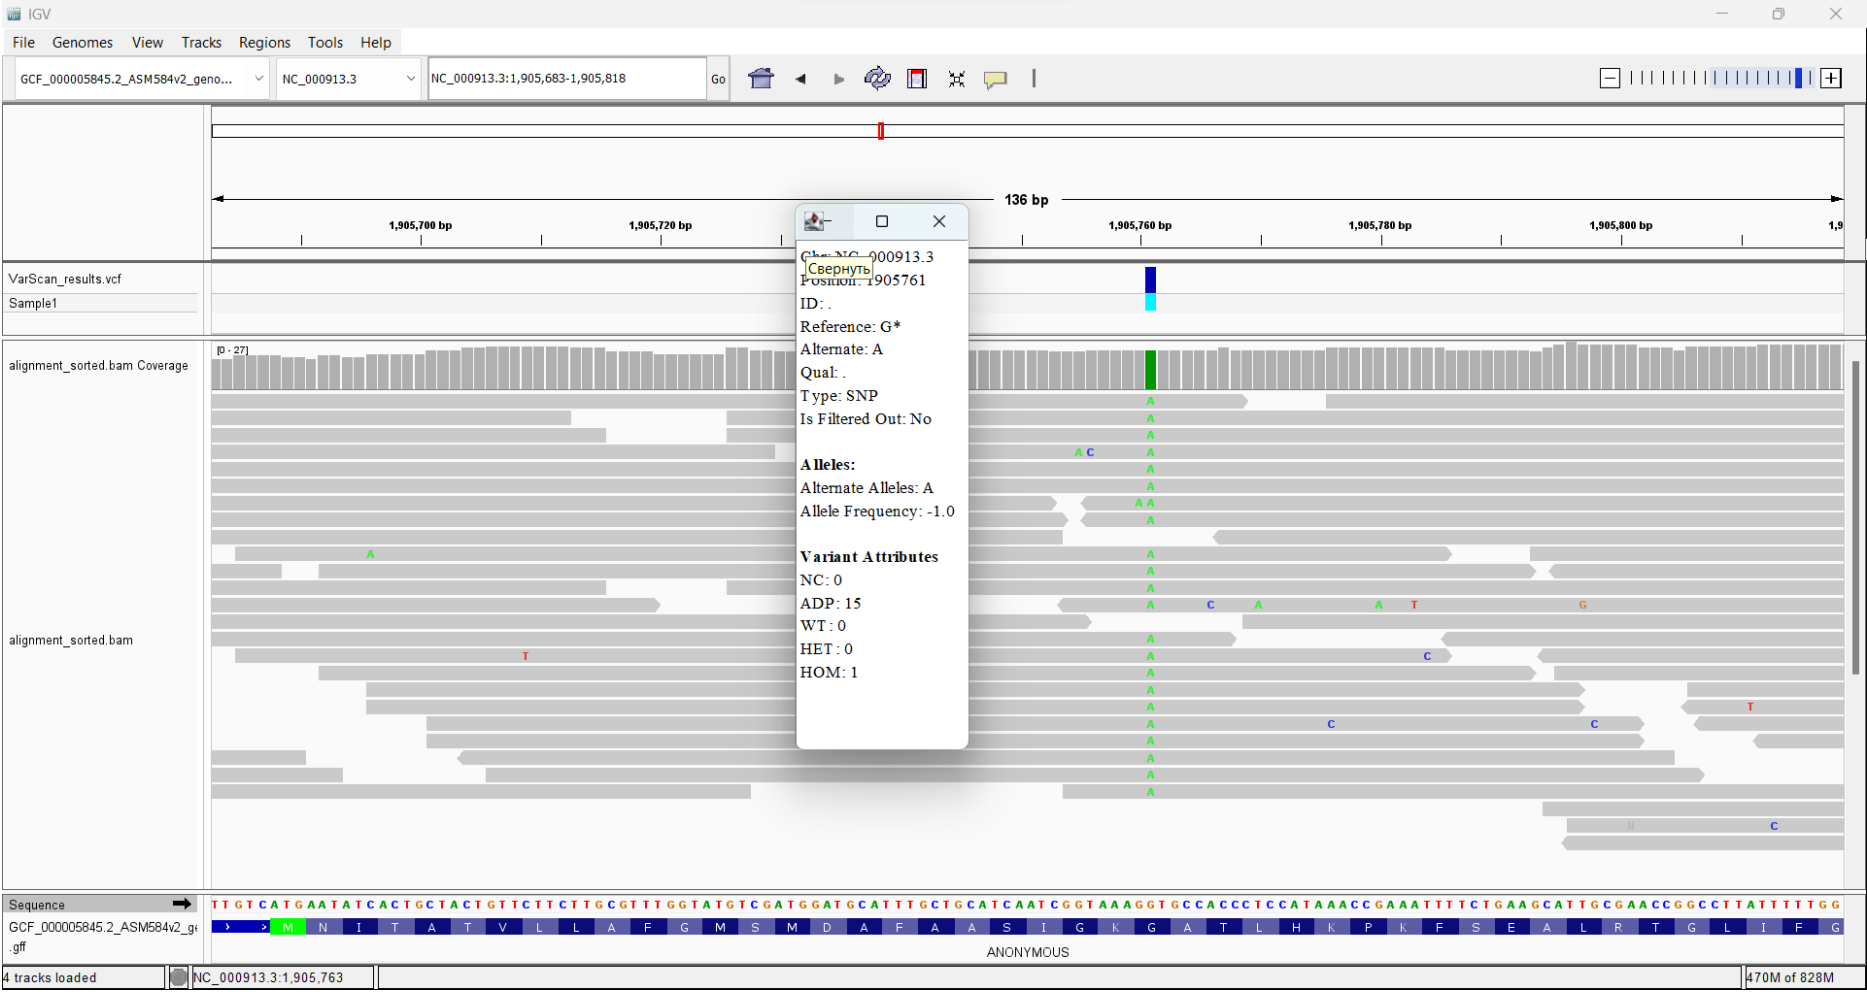
SNP 3

- Position: 852762
- Reference: A
- Alternate: G
- Gene: rybA
- Effect: Upstream gene variant with no amino acid change.



SNP 4

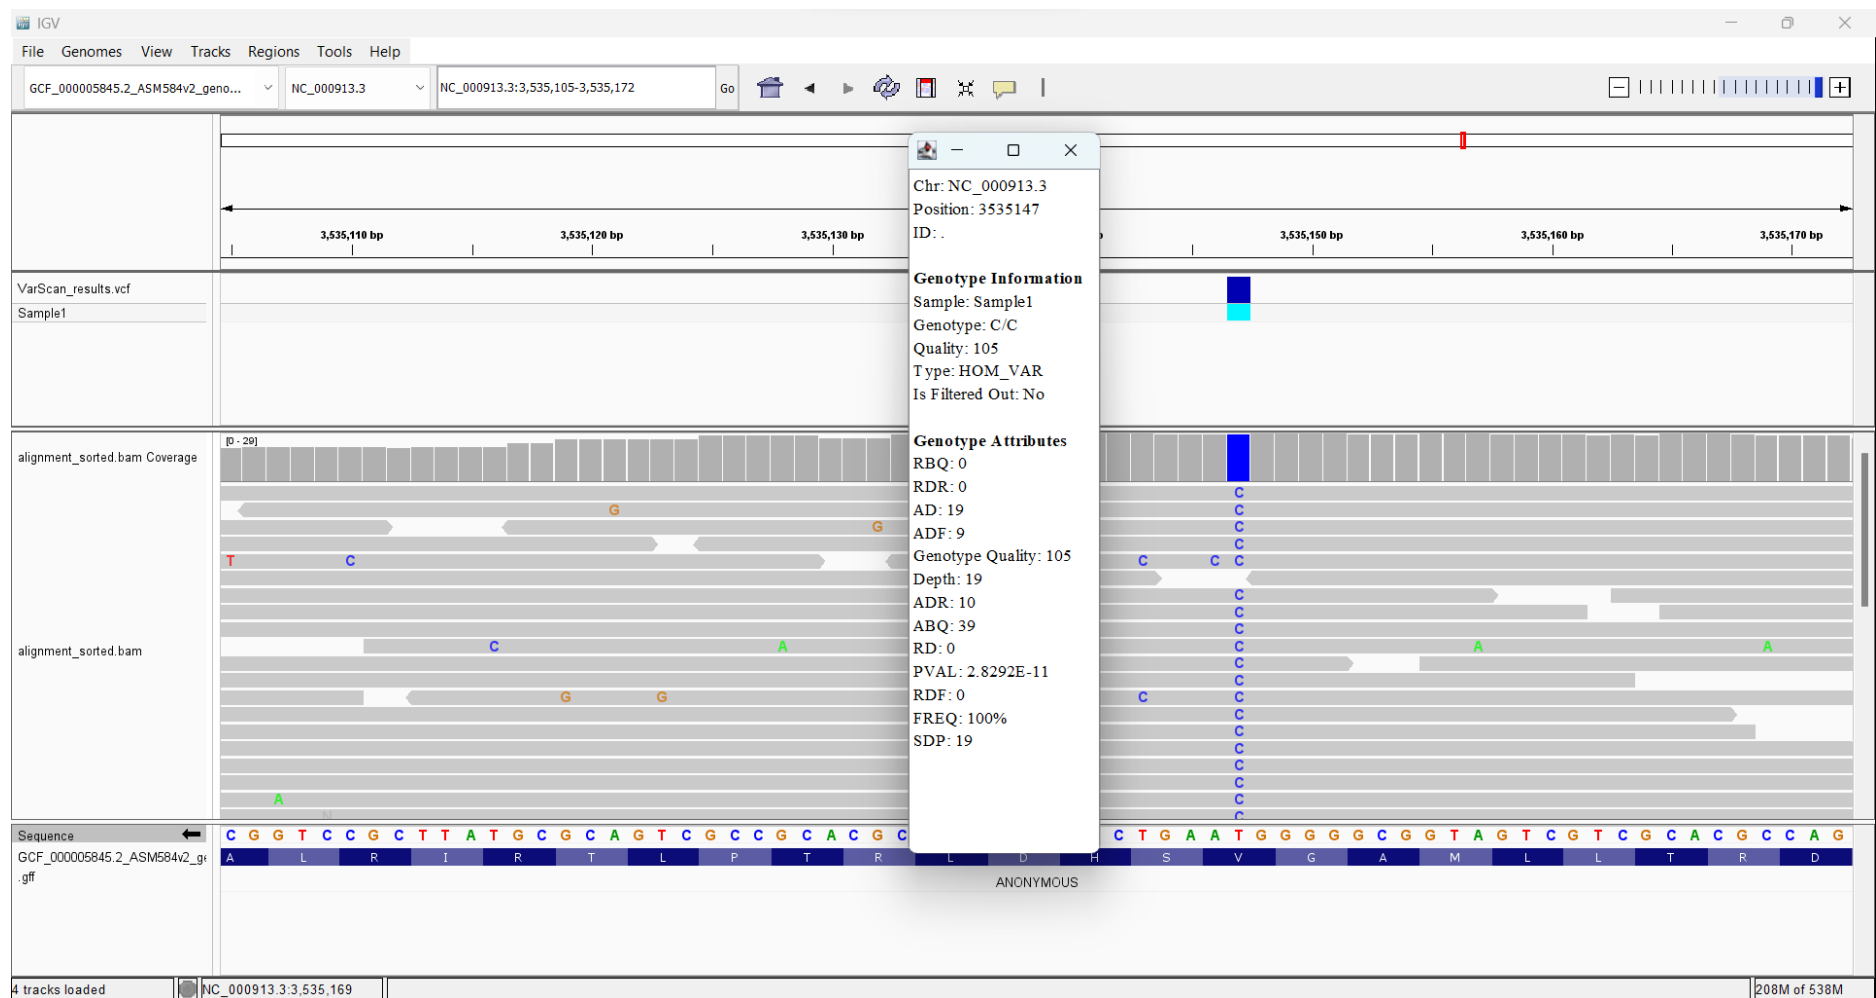
- Position: 1905761
- Reference: G
- Alternate: A
- Gene: mntP
- Effect: Missense variant, changing amino acid from Glycine (Gly) to Aspartic Acid (Asp)



SNP 5

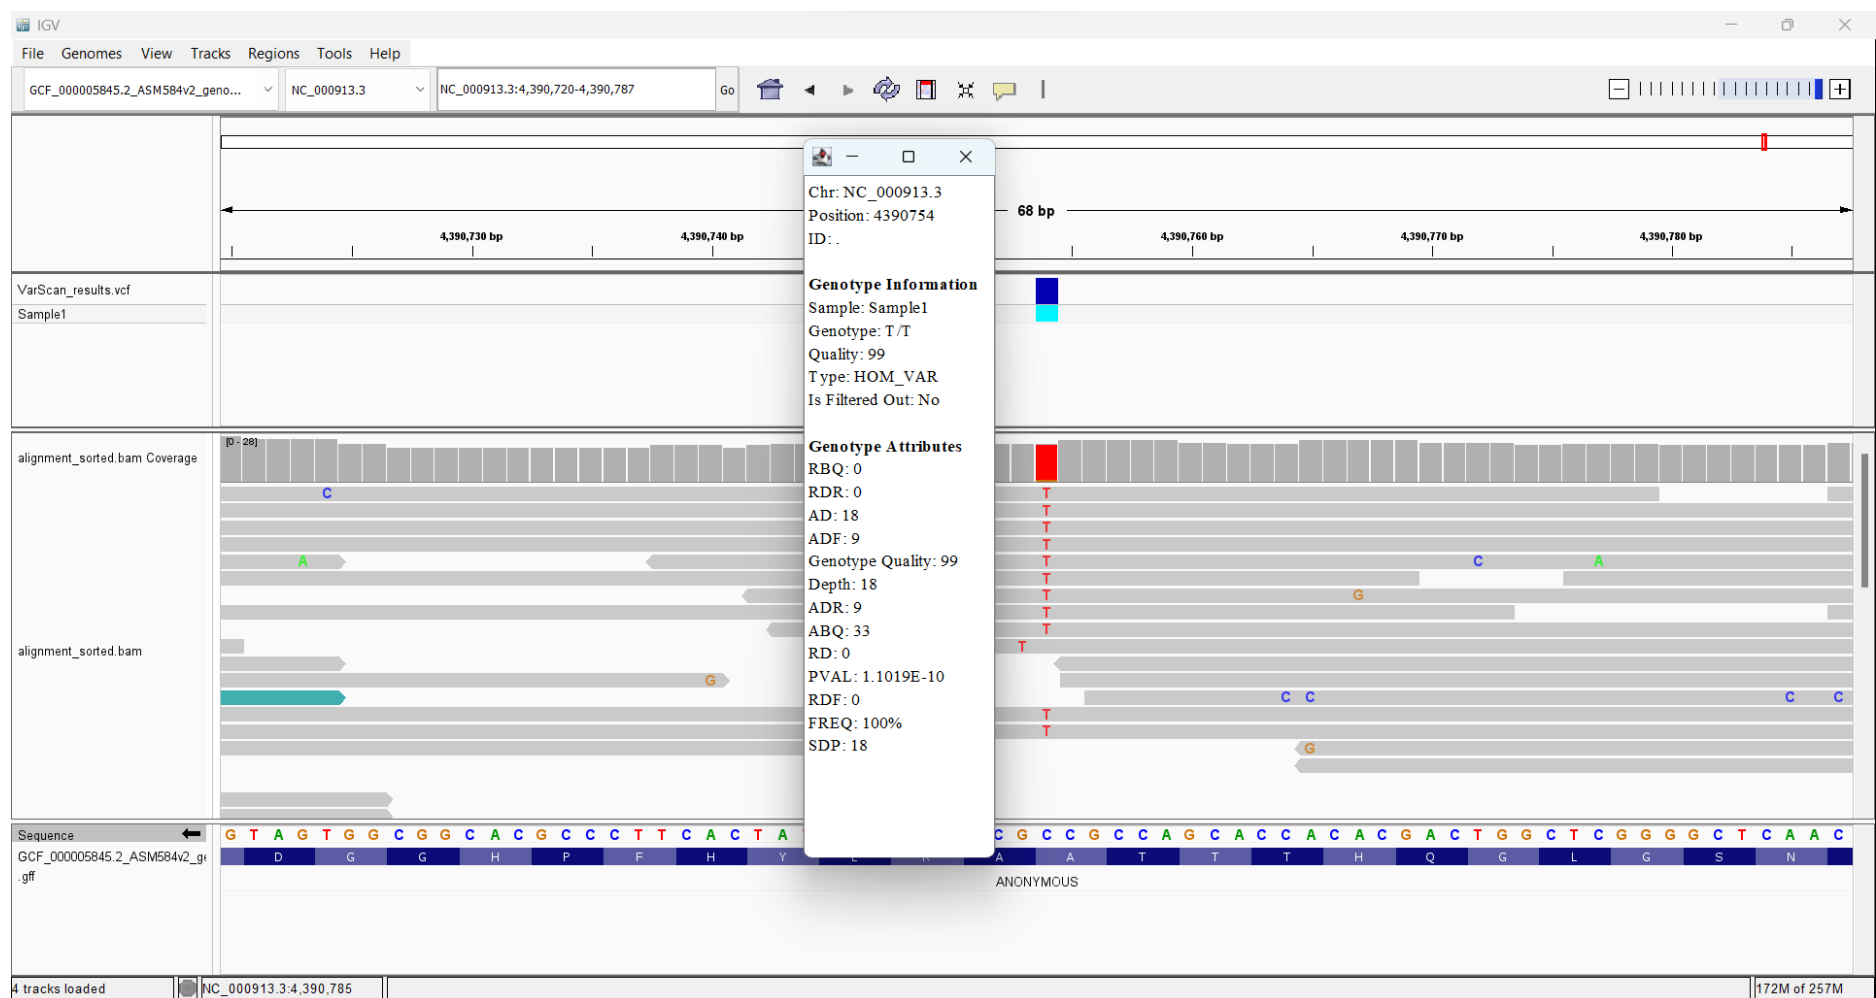
- Position: 3535147
- Reference: T
- Alternate: G
- Gene: envZ

- Effect: Missense variant, changing amino acid from Valine (Val) to Glycine (Gly)



SNP 6

- Position: 4390754
- Reference: C
- Alternate: A
- Gene: rsgA
- Effect: Synonymous variant, with no change in the amino acid sequence Ala → Ala



8. Automatic SNP annotation

```
conda install -c bioconda snpeff
```

```
wget https://ftp.ncbi.nlm.nih.gov/genomes/all/GCF/000/005/845/GCF_000005845.2_ASM584v2/GCF_000005845.2_ASM584v2
```


`_genomic.gbff.gz`

Create empty text file `snpEff.config`, and add there just one string: `k12.genome : ecoli_K12`

Create folder for the database

Put there your `.gbk` file (unzip and rename to `genes.gbk`)

Create database

Annotate

```
echo "k12.genome : ecoli_K12" > snpEff.config
mkdir -p data/k12
gunzip GCF_000005845.2_ASM584v2_genomic.gbff.gz
cp GCF_000005845.2_ASM584v2_genomic.gbff data/k12/genes.gbk
snpEff build -genbank -v k12
snpEff ann k12 VarScan_results.vcf > VarScan_results_annotated.vcf
```

As a result, there will be a `vcf` file with additional field "ANN" (for "annotation"), describing all the effects for each SNP.

[VarScan_results_annotated.vcf](#)

Results from automatic annotation correlate with acquired with VarScan